Data Preparation For Data Science

MINKAI WU

Outline

Data Acumen

- Data Science Process
- Data Quality
- Data Source
- Data File Format
- Data Types

Data Cleaning

- Missing Data
- Invalid Data
- Feature extraction
- Demo

Web Data Preparation

- Understanding the HTML Page Structure
- Python and Regular Expressions to clean data
- Python and Beautiful Soup to collect data
- Demo

Introduction: Data Science Process

1. Problem Statement
2. Data Collection & Storage
3. **Data Preparation**
    1. **Access Data**
    2. **Clean Data**
    3. **Transform Data**
4. Data Analysis & Visualization
5. Modeling
6. Presentation or Productize

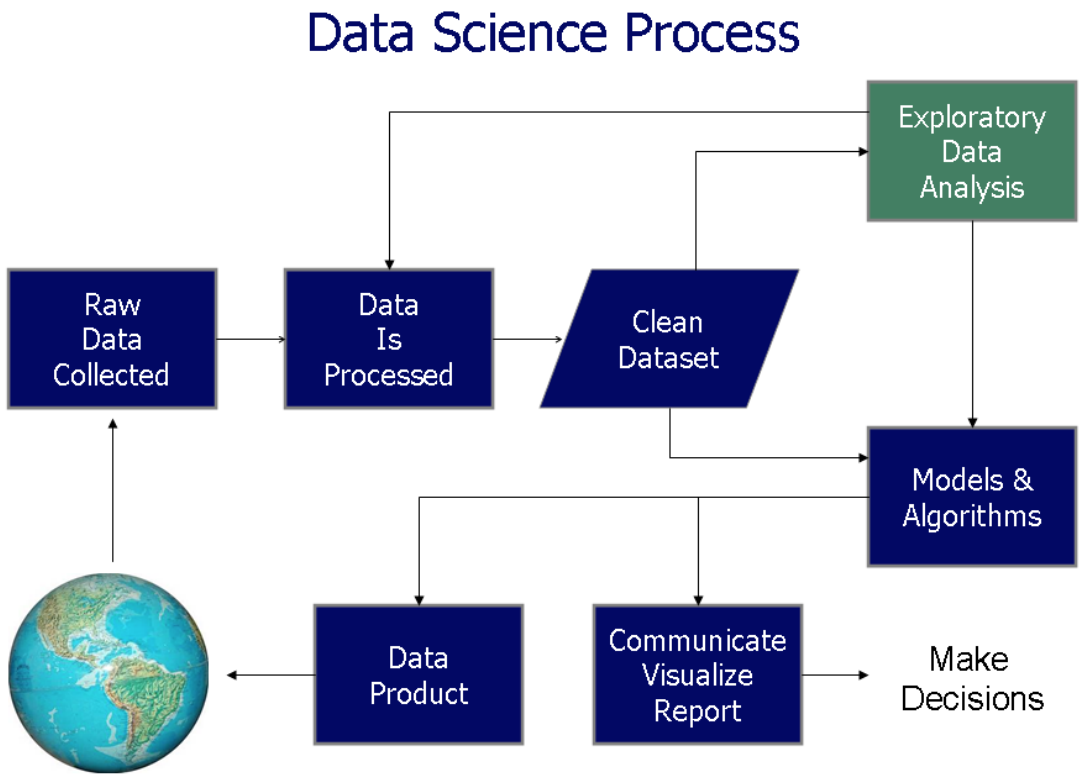

Introduction: Data Quality Issues

- Incorrect/Invalid Entry
  - age = 203; gender = X; price = -100; weekday=8
- Missing Data
  - N/A; Null; ” “; Unknown
- Unstructured Data
  - merged cell; double header; html
- Conflicting Data
  - revenue =1000; unit = 0
- Duplicates
  - double loading; double counting
- Outlier
  - House Price = $1B

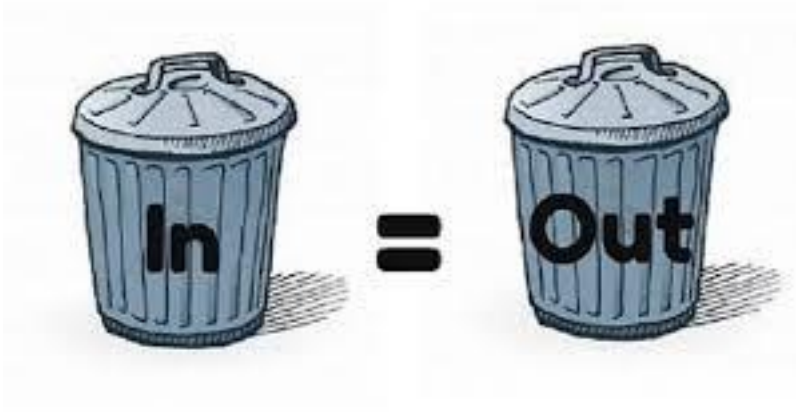

Introduction: Data Source

- Data File
- Database/Data Warehouse
- Web Data
- Big Data Platform

Data File: Structured Data

- Excel:
  - Most common; most problematic
- Delimited format
  - Most common; most preferred
  - Common delimited (csv); tab delimited(tsv); “|” delimited
  - Problem: delimiter in data field. E.g. Los Angles, CA
  - Problem: encoding
- Fixed length
  - Every column has fixed length
  - Problem: Oversized column

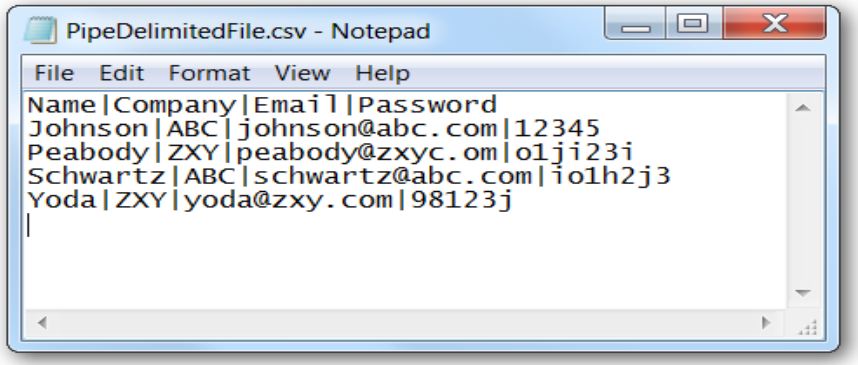

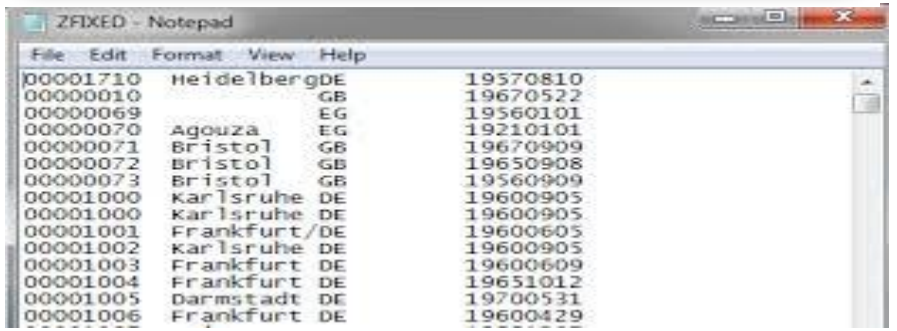

Data File: JSON

- JavaScript Object Notation
- Semi- Structured
- Attributes are on the left-hand side of colon
- Values are on the right-hand side of colon
- Attributes are separated by a comma
- Multi-value attributes are as hierarchical values

```
{
  "firstName": "Sally",
  "birthDate": "1971-09-16",
  "faveColor": "light\"Carolina\" blue",
  "pet":
  [
    {
      "type": "dog",
      "name": "Fido"

    },
    {
      "type": "dog",
      "name": "Lucky"
    }
  ],
  "job": {
    "jobTitle": "Data Scientist",
    "company": "Data Wizards, Inc.",
    "salary":129000
  }
}
```

Data File: XML and Parquet

- XML
  - Extensible Markup Language
  - Semi-Structured
  - Most common for data exchange

```
<?xml version="1.0" standalone="no"?>
<GridView>
    <rowheader>
        <colheader text="FirstName" width="80" />
        <colheader text="LastName" width="80" />
        <colheader text="Company" width="120" />
        <colheader text="E-mail" width="160" />
    </rowheader>
    <row>
        <col text=" " backcolor="-1" forecolor="-16777216" />
        <col text=" " backcolor="-1" forecolor="-16777216" />
        <col text=" " backcolor="-1" forecolor="-16777216" />
        <col text=" " backcolor="-1" forecolor="-16777216" />
    </row>
    <row>
        <col text="John" backcolor="-1" forecolor="-16777216" />
        <col text="Doe" backcolor="-1" forecolor="-16777216" />
        <col text="Microsoft" backcolor="-7722014" forecolor="-32944" />
        <col text="joe@aol.com" backcolor="-1" forecolor="-16777216" />
    </row>
```

- Parquet
  - Column Store
  - Spark

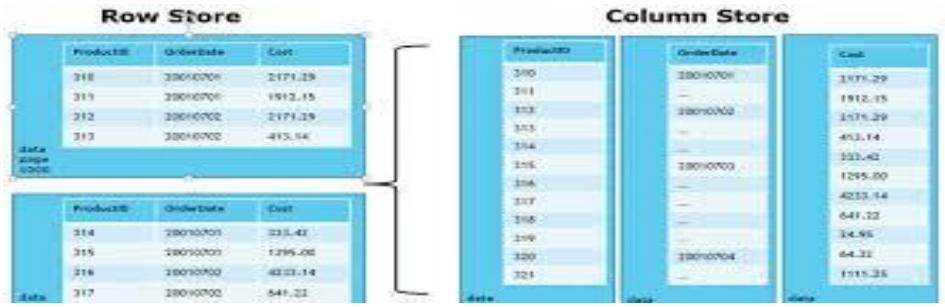

Web Data: HTML - Unstructured

Secure | https://www.indeed.com/m/jobs?q=data+scientist&l=Los+Angeles%2C+CA

indeed

data scientist jobs in Los Angeles, CA

Jobs 1-10 of 555: **All** - New - Be the first to see new jobs

Data Scientist
The Honest Company - Los Angeles, CA
Desired Experience: Hive, Machine Learning, R, C/C++, MATLAB, Data Mining, Scala, Weka, Java, Spark, Python
1 day ago

Data Scientist
Fuel Cycle - Los Angeles, CA
$120,000 - $160,000 a year
Desired Experience: Machine Learning, R, MySQL, AI, Data Mining, Sas, Java, Data Science, Python
30+ days ago

Data Scientist/Quantitative Analyst
Magid - Los Angeles, CA
Desired Experience: Machine Learning, R, Git
8 days ago

Data Scientist
Kaiser Permanente - Pasadena, CA
20 days ago

Data Scientist
L.A. Care Health Plan - Los Angeles, CA 90017
Desired Experience: Machine Learning, R, Sas, Tableau, Spark, Data Science, Python
11 days ago

Senior Data Scientist
Ticketmaster - Hollywood, CA 90028
Desired Experience: Machine Learning, C/C++, Hadoop, HBase, Java, Spark, Python
15 hours ago

```
<!DOCTYPE html PUBLIC "-//WAPFORUM//DTD XHTML Mobile 1.0//EN" "http://www.wapforum.org/DTD/xhtml-mobile10.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
<head>
<title>Data Scientist Jobs, Employment in Los Angeles, CA | Indeed Mobile</title>
<meta name="description" content="555 Data Scientist Jobs available in Los Angeles, CA on Indeed.com.  one search.  all jobs.">
<meta http-equiv="content-type" content="text/html; charset=utf-8" />
<meta name="referrer" content="origin">
<link rel="next" href="jobs?q=data+scientist&l=Los+Angeles%2C+CA&start=10&pp=AAoAAAFesdLhCAAAAAEdI4JuAQACnhKpqjRrCB4PDNNsJljfj2tIIH6kng">
<link rel="canonical" href="/q-Data-Scientist-l-Los-Angeles,-CA-jobs.html"/>
<link rel="alternate" href="android-app://com.indeed.android.jobsearch/https/www.indeed.com/m/jobs?l=Los+Angeles%2C+CA&q=data+scientist">
<style type="text/css"><!--


--></style>
</head>
<body>
<div><a href="/m/"><img src="/m/_xhtmlmp/images/indeed_r.gif" height="26" width="99" alt="Indeed Mobile Job Search"/></a></div>
<h1 class="serpHeading">data scientist jobs in Los Angeles, CA</h1>
<p>Jobs 1-10 of 555:
<b>All</b> - <a rel="nofollow" href="/m/jobs?q=data+scientist&l=Los+Angeles%2C+CA&from=newbtn&fromage=last">New</a>
 - <a href="/m/jobalerts?q=data+scientist&l=Los+Angeles%2C+CA&dest=%2Fm%2Fjobs%3Fq%3Ddata%2Bscientist%26l%3DLos%2BAngeles%252C%2BCA">Be the first to see
 new jobs</a>
</p>
<p><h2 class="jobTitle"><a rel="nofollow" href="viewjob?jk=46caf455b09ff764">Data Scientist</a></h2><br/>
The Honest Company - <span class="location">Los Angeles, CA</span><br/>
<style type="text/css"><!--
/* desired experience */
.experienceHeader{color:#666}
.experienceList{color:#000}
.experience{margin-top: 4px; margin-bottom: 5px}

--></style>
<div class="experience">
<span class="experienceHeader">Desired Experience: </span><span class="experienceList">Hive, Machine Learning, R, C/C++, MATLAB, Data Mining, Scala,
Weka, Java, Spark, Python</span><br>
</div>
<span class="date">1 day ago</span></p>
<hr/>
<p><h2 class="jobTitle"><a rel="nofollow" href="viewjob?jk=7b8f1e2c8b577bf6">Data Scientist</a></h2><br/>
Fuel Cycle - <span class="location">Los Angeles, CA</span><br/>
<span class="salary">$120,000 - $160,000 a year</span><br/><style type="text/css"><!--
```

Data Source: RDBMS


```python
import MySQLdb

# Open database connection
db = MySQLdb.connect("localhost","testuser","test123","TESTDB" )

# prepare a cursor object using cursor() method
cursor = db.cursor()

# execute SQL query using execute() method.
cursor.execute("SELECT VERSION()")

# Fetch a single row using fetchone() method.
data = cursor.fetchone()

print "Database version : %s " % data

# disconnect from server
db.close()
```

Data Source: Big Data Platform

- HDFS – Hive
  - Text file and table
- Spark – RDD
  - Resilient Distributed Datasets
  - RDD is a read-only, partitioned collection of records
- Amazon -- S3
  - Cloud Data storage
  - File can be in any format

Data Types

- Numeric
  - Discrete: Count; Rating; Grade; Fibonacci Series
  - Continuous: Revenue; Distance; Home Value
  - Watch out: data range!
- Binary (Dummy)
  - Special case of numeric
  - E.g.: IsMale; HasCar; Pass
- Categorical
  - Usually contains characters: Gender, Product, Geo, etc.
  - Can be consist of pure numbers: SSN, Zipcode, Phone Number
  - Watch out: Valid Values
- Dates and Time
  - Date, Time, Datetime, Timestamp
  - Watch out: Time Zone! UTC=Coordinated Universal Time = GMT = Greenwich Mean Time
- Missing

Data Types: Missing

- Null
  - Absence of everything; missing; empty
- Blank
  - “ “ or “ “ or any invisible characters
  - can mean missing
  - can mean “N/A”
- N/A
  - Can mean “not available”: e.g. Age
  - Can mean “not applicable”: e.g. Middle Name
  - Can mean “no answer”: e.g. Customer Satisfaction Rating on a Questionnaire

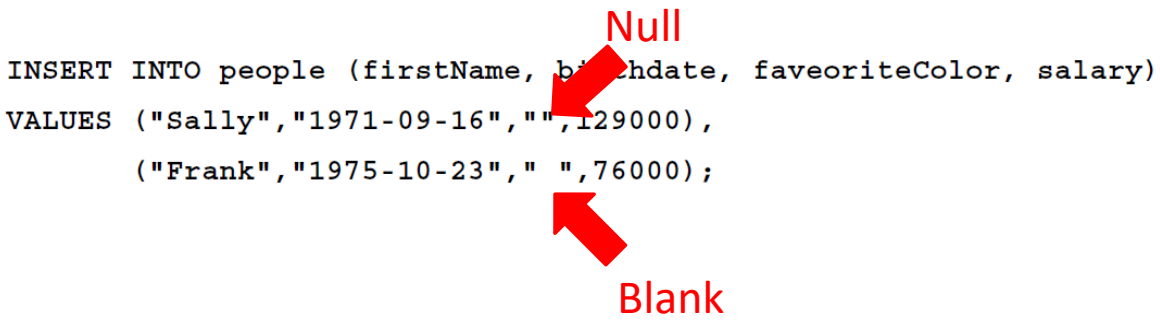

Data Preparation Best Practice

Data Preparation Steps

- **Data Cleansing**
  - **Integrate (mapping)**: integrate various data sources into one dataset. E.g. sales units, sales revenue, price
  - **Conform**: Conform the inconsistent values. E.g. Na, n/a => missing; Los Angeles, L.A. => LA
  - **Filter**: Filter out the columns and rows not needed for modeling
  - **Extract**: Extract new column/feature from existing columns. E.g. month from date
  - **Group**: Group many categorical values into less buckets
  - **Aggregate**: Aggregate/Disaggregate date to the desired granularity
  - **Derived feature**: Calculate new metrics based on existing metrics. E.g. Price =Revenue/Units
- **Handle Missing Data**
- **Identity Outlier**
- **Transform Data**
  - **One hot encoding: categorical to numerical**
  - **Normalization/Standardization**
  - **Log transformation**

Data Cleansing: Regex 101

a single character of: a, b or c	**[abc]**
a character except: a, b or c	**[^abc]**
a character in the range: a-z	**[a-z]**
a character not in the range: a-z	**[^a-z]**
a character in the range: a-z or A-Z	**[a-zA-Z]**
any single character	**.**
any whitespace character	**\s**
any non-whitespace character	**\S**
any digit	**\d**
any non-digit	**\D**
any word character	**\w**
any non-word character	**\W**

capture everything enclosed	**(...)**
match either a or b	**(a|b)**
zero or one of a	**a?**
zero or more of a	**a***
one or more of a	**a+**
exactly 3 of a	**a{3}**
3 or more of a	**a{3,}**
between 3 and 6 of a	**a{3,6}**
start of string	**^**
end of string	**$**

Data Cleansing: Useful Regex

- Replace
  - Reverse last name and first name: San, Zhang => Zhang San
  - Regex=/([a-zA-Z]+),\s*([a-zA-Z]+)/, Replace = $2 $1
- Extract
  - Extract url from html: amgheziName
  - Regex = /href=/"([^"]*)/, Replace = $1
- Validation
  - Validate a valid email
  - Regex =/^([a-z0-9_\.-]+)@([\da-z\.-]+)\.([a-z\.]{2,6})$/i

Missing Data: Types

- **Missing completely at random: MCAR**
  - Roll a dice
  - Lottery number
- **Not missing at random: NMAR**
  - missing values are systematic
  - Income: higher income is less likely to respond
  - Weight: higher weight is less likely to respond
  - Smoking
- **Missing at random: MAR**
  - Most Common
  - Missing values can somewhat be predicted by known info
  - Know height, missing weight
  - Know # of rooms, missing sqrt

Missing Data: Handling

- Impute from other attributes
- Impute from other observations
  - Majority vote (categorical)
  - Mean of same/similar group (numerical)
  - Carry last value (time series)
  - Linear fill (time series)
  - Carry same trend (time series)
- “Missing” Category (not missing at random)
- Extra indicator
- Logical estimation
- Remove row or column

Outliers: 1.5 IQR

Check the frequency distribution of the data

Box-plot: An outlier is a point of data that lies over 1.5 IQRs below the first quartile (Q1) or above third quartile (Q3) in a given data set.

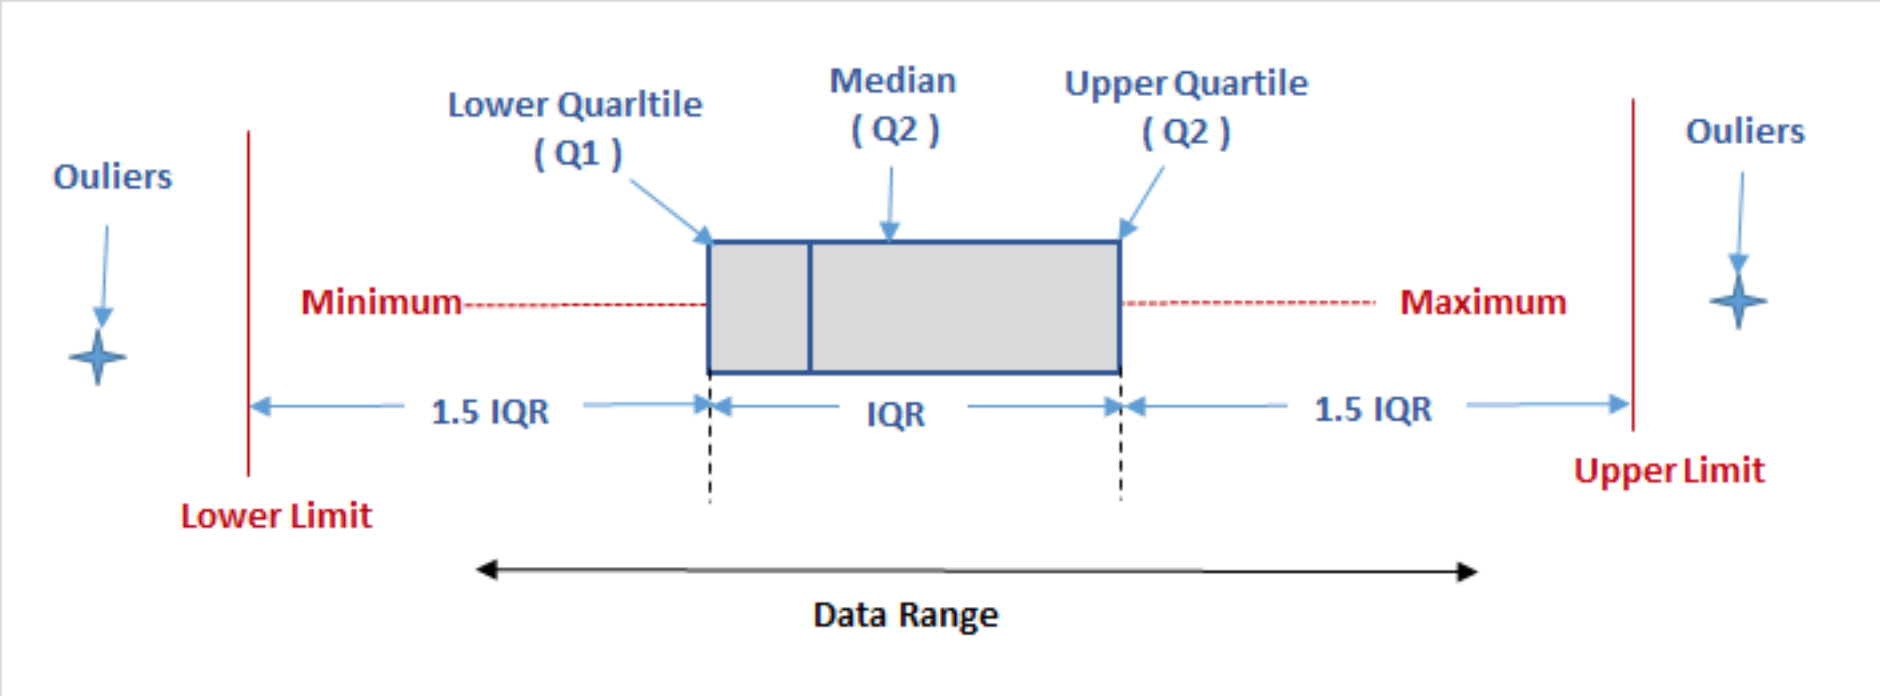

Outlier: Normal Distribution

Outlier: 2 or 3 STD from mean

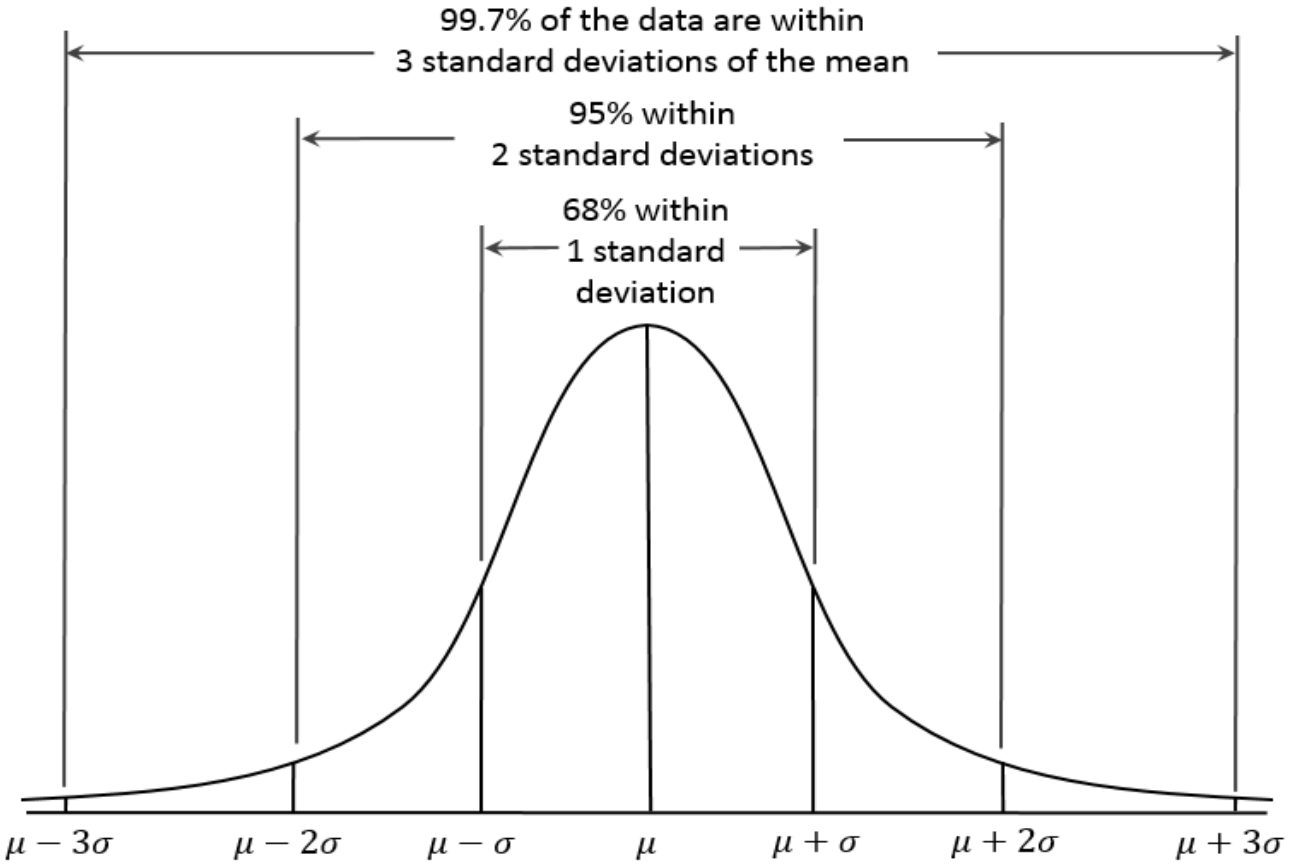

Outlier: Other Technics

- Univariable Outlier:
  - Median Absolute Deviation
- Multivariate Outlier
  - Mahalanobis Distance

Data Transformation: Normalization vs Standardization

	Normalization	Standardization
Formula	$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$	$x_{new} = \frac{x - \mu}{\sigma}$
Pro	• Bounded (-1,1) • Apply to all distribution	• Works well for normal distribution
Con	• Make outliers “normal”	• Unbounded • Only works well for normal distribution

Transformation: When to Normalize

- Linear Model
  - Recommended
  - Doesn’t change model accuracy
  - Easier to compare coefficient: larger coefficient, larger impact
  - Intercept well interpreted: the expected value of Yi when the predictors are set to their means
  - Avoid coefficient like 10^-9 when one variable has a very large scale
  - More difficult to interpret the model in terms of on unit change in Xi
- Tree Model
  - Not necessary as the scale is irrelevant
- Logistic Regression
  - Typically not needed
- SVM
  - Recommended
  - Help with faster converge

Transformation: Log

Linear Model; Skewed Data

- Log Predictor

$$y = e^{ax} + b \xrightarrow{\log x} y = ax' + b$$

- Log Outcome

$$y = ln\,(ax + b) \xrightarrow{\log y} y' = ax + b$$

- Log both

$$y = e^{c} * {x_1}^{a} * {x_2}^{b} \xrightarrow{\text{yields}} \ln y = c + ax_1 + bx_2$$

Demo

Use Python to clean Airbnb listings data (from file)

Web Data Preparation

WEB data raw format: HTML

Understanding the HTML Page Structure

HTML can be parsed in two ways:

- The line-by-line delimiter model
- The tree structure model

```
<div id="content">
<h2>Sep 13, 2014</h2>

<a href="/2014/sep/14/">← next day</a> Sep 13, 2014  <a
  href="/2014/sep/12/">previous day →</a>

<ul id="ll">
<li class="le" rel="petisnnake"><a href="#1574618"
  name="1574618">#</a> <span style="color:#b78a0f;8"
  class="username" rel="petisnnake">&lt;petisnnake&gt;</span> i
  didnt know that </li>
...
</ul>
...
</div>
```

Web Scraping: Line by Line

The line-by-line delimiter model

```
<div id="content">
<h2>Sep 13, 2014</h2>

<a href="/2014/sep/14/">← next day</a> Sep 13, 2014  <a
  href="/2014/sep/12/">previous day →</a>

<ul id="ll">
<li class="le" rel="petisnnake"><a href="#1574618"
  name="1574618">#</a> <span style="color:#b78a0f;8"
  class="username" rel="petisnnake">&lt;petisnnake&gt;</span> i
  didnt know that </li>
...
</ul>
...
</div>
```

- <h2></h2> tags as delimiters to extract the date
- tags as delimiters to extract text
- Rel=“” as delimiters to extract user name
- From the end of to the beginning of is the actual line message

Extract date by Regex: <h2>(.+)<\/h2>

Extract message by Regex : <\/span>(.+)<\/li>

Web Scraping: Tree Model

The tree structure model: we can consider the structure of HTML as a tree structure

```
<div id="content">
<h2>Sep 13, 2014</h2>

<a href="/2014/sep/14/">← next day</a> Sep 13, 2014  <a
  href="/2014/sep/12/">previous day →</a>

<ul id="ll">
<li class="le" rel="petisnnake"><a href="#1574618"
  name="1574618">#</a> <span style="color:#b78a0f;8"
  class="username" rel="petisnnake">&lt;petisnnake&gt;</span> i
  didnt know that </li>
...
</ul>
...
</div>
```

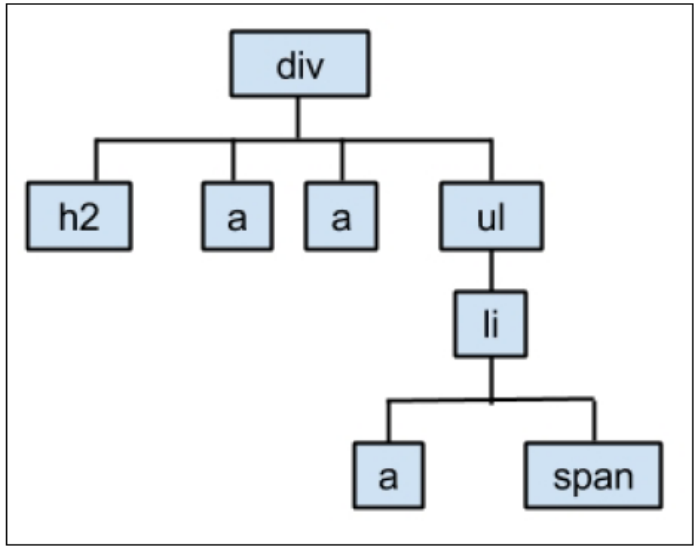

Extract date by beautifulSoup: div.h2.text

Extract message by beautifulSoup: div.ul.li.text

Demo

Use Python Beautifulsoup to collect and clean job listing data from indeed.com