




Artificial Neuro Networks (ANN) Classifier


WEEK 3

Neuron



Perceptron

- A classifier that maps input vector to a single binary output.
- Example Illustration:




- Commonly used example: A sign function of a linear combination of the input:

$$y = \mathcal{F} \left(\sum_{i=1}^N w_i x_i + \theta \right), \quad \mathcal{F}(s) = \begin{cases} 1 & \text{if } s > 0 \\ -1 & \text{otherwise.} \end{cases}$$


Feedforward Neural Network & Multi-layer Perceptron

- Feedforward Neural Network:
 - Layered Structure.
 - Each node receives inputs from nodes directly below, and outputs to the nodes above them.
 - No connection within a layer.
- Multi-layer perceptron is a common practice to implement feedforward neural network.
- In the slides we focus on feedforward neural networks.



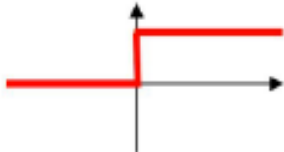
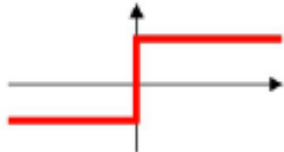


Multi-layer Perceptron and Neural Network

- Intuition: we can approximate almost any 1-D functions with a combination of step functions.




Perceptron

Example Activation Function:


Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	

Neural Network

- Choice of Activation Function
 - Sign function, step function, piecewise linear function are hard to deal with.
 - Either non-continuous, or their derivatives are not continuous.
 - Commonly use a “**sigmoid**”
 - A logistic function.
 - $\sigma(z) = \frac{1}{1 + e^{-z}}$
 - First order derivative: $\frac{d\sigma}{dz} = (1 - \sigma)\sigma$




(b) Sign function



(c) Sigmoid function

Neural Network

- Nice properties by using multiple output nodes in a neural network:
 - **Regression**
 - Multi-dimensional regression
 - **Multi-class Classification:**
 - Use one-hot encoding in class labels
 - Eg. Class 1: $[1, 0, 0, 0, \dots, 0]$
Class 2: $[0, 1, 0, 0, \dots, 0]$
 - **Joint Binary Predictions:**
 - Image tagging



Notations


x_i : feature i of input variable x .

\hat{y}_j : j -th output variable.


y_j : j -th true output label.

$w_{ij}^{(k)}$: weight of the output of node j in layer k to node i in layer $k + 1$.

$h_i^{(k)}$: output of node i of hidden layer k .



Forward Propagation and Backpropagation



- Forward pathing to make predictions and obtain prediction error.
- Backward pathing to learn the model parameters using the prediction error.

Forward Propagation


- Forward pathing to make predictions

$$h_j = \sigma\left(\sum_i x_i w_{ji}^{(1)}\right)$$

$$\hat{y}_k = \sigma\left(\sum_j h_j w_{kj}^{(2)}\right)$$

- Then calculate the error, or the cost function:

$$\text{cost function: } J(w, x, y) = \frac{1}{2} \sum_k (\hat{y}_k - y_k)^2$$



Backpropagation


- Backpropagation is short for “backward propagation of errors”.
- Commonly referred to the **gradient descent** method to find the optimal weights.
- Recall the general gradient descent method:

Repeat until convergence:

{

$$w_i \leftarrow w_i - \alpha \frac{\partial J(w)}{\partial w_i}, \text{ for every } i$$

}




Backpropagation

Recall: Chain rule for partial derivatives:

Let $x = x(u, v)$ and $y = y(u, v)$ have first-order partial derivatives at the point (u, v) and suppose that $z = f(x, y)$ is differentiable at the point $(x(u, v), y(u, v))$. Then $f(x(u, v), y(u, v))$ has first-order partial derivatives at (u, v) given by

$$\begin{aligned}\frac{\partial z}{\partial u} &= \frac{\partial z}{\partial x} \frac{\partial x}{\partial u} + \frac{\partial z}{\partial y} \frac{\partial y}{\partial u} \\ \frac{\partial z}{\partial v} &= \frac{\partial z}{\partial x} \frac{\partial x}{\partial v} + \frac{\partial z}{\partial y} \frac{\partial y}{\partial v}.\end{aligned}$$

Backpropagation



$$h_j = \sigma\left(\sum_i x_i w_{ji}^{(1)}\right)$$

$$\hat{y}_k = \sigma\left(\sum_j h_j w_{kj}^{(2)}\right)$$

$$J(\underline{w}, \underline{x}, \underline{y}) = \frac{1}{2} \sum_k (\hat{y}_k - y_k)^2$$


Calculate the gradient

$$\begin{aligned} \frac{\partial J(w)}{\partial w_{kj}^{(2)}} &= \frac{\partial \frac{1}{2} \sum_k (\hat{y}_k - y_k)^2}{\partial w_{kj}^{(2)}} \\ &= \frac{\partial \frac{1}{2} \sum_k (\hat{y}_k - y_k)^2}{\partial \hat{y}_k} \cdot \frac{\partial \hat{y}_k}{\partial w_{kj}^{(2)}} \\ &= (\hat{y}_k - y_k) \frac{\partial \sigma\left(\sum_j h_j w_{kj}^{(2)}\right)}{\partial w_{kj}^{(2)}} \\ &= (\hat{y}_k - y_k) \sigma' \left(\sum_j h_j w_{kj}^{(2)} \right) h_j \end{aligned}$$

Where $\sigma'(z) = \sigma(z)(1 - \sigma(z))$

Similar gradient can be calculated for $\frac{\partial J(w)}{\partial w_{ji}^{(1)}}$.

Backpropagation (Cont'd)




$$\frac{\partial J(w)}{\partial w_{kj}^{(2)}} = (\hat{y}_k - y_k) \sigma' \left(\sum_j h_j w_{kj}^{(2)} \right) h_j$$

$$\frac{\partial J(w)}{\partial w_{ji}^{(1)}} = \sum_k (\hat{y}_k - y_k) \sigma' \left(\sum_j h_j w_{kj}^{(2)} \right) w_{kj}^{(2)} \cdot \sigma' \left(\sum_i w_{ji}^{(1)} x_i \right) x_i$$

Computational cost is linear
with feedforward NN depth!

Dropout



First Example

Inputs			Output
0	0	1	0
1	1	1	1
1	0	1	1
0	1	1	0

Denotation

Variable	Definition
X	Input dataset matrix where each row is a training example
y	Output dataset matrix where each row is a training example
I_0	First Layer of the Network, specified by the input data
I_1	Second Layer of the Network, otherwise known as the hidden layer
syn_0	First layer of weights, Synapse 0, connecting I_0 to I_1 .
$*$	Elementwise multiplication, so two vectors of equal size are multiplying corresponding values 1-to-1 to generate a final vector of identical size.
$-$	Elementwise subtraction, so two vectors of equal size are subtracting corresponding values 1-to-1 to generate a final vector of identical size.
$x.dot(y)$	If x and y are vectors, this is a dot product. If both are matrices, it's a matrix-matrix multiplication. If only one is a matrix, then it's vector matrix multiplication.

Second Example

Inputs			Output
0	0	1	0
0	1	1	1
1	0	1	1
1	1	1	0

Denotation

Variable	Definition
X	Input dataset matrix where each row is a training example
y	Output dataset matrix where each row is a training example
I_0	First Layer of the Network, specified by the input data
I_1	Second Layer of the Network, otherwise known as the hidden layer
I_2	Final Layer of the Network, which is our hypothesis, and should approximate the correct answer as we train.
syn_0	First layer of weights, Synapse 0, connecting I_0 to I_1 .
syn_1	Second layer of weights, Synapse 1 connecting I_1 to I_2 .
I_2_error	This is the amount that the neural network "missed".
I_2_delta	This is the error of the network scaled by the confidence. It's almost identical to the error except that very confident errors are muted.
I_1_error	Weighting I_2_delta by the weights in syn_1 , we can calculate the error in the middle/hidden layer.
I_1_delta	This is the I_1 error of the network scaled by the confidence. Again, it's almost identical to the I_1_error except that confident errors are muted.